

3_loan_prediction

吴清柳

July 24, 2023

1 Loan Prediction with Logistic Regression

使用线性回归模型的贷款预测

	precision	recall	f1-score	support
0	0.92	0.43	0.59	51
1	0.82	0.99	0.89	134
accuracy			0.83	185
macro avg	0.87	0.71	0.74	185
weighted avg	0.85	0.83	0.81	185

accuracy is 0.8324324324324325

```
[2]: import os
import numpy as np
import pandas as pd
import warnings

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score

from sklearn.linear_model import LogisticRegression

# default theme
sns.set(context='notebook', style='darkgrid', palette='deep', font='sans-serif', font_scale=1, color

# warning handle
warnings.filterwarnings('ignore')
```

```
[12]: # set train and test file path
tr_path=''
te_path=''

for dirname, _, filenames in os.walk('../dataset/'):
    for filename in filenames:
        cur_path=os.path.join(dirname,filename)
        print(cur_path)
        if filename.find('train')==0:
            tr_path=cur_path
        elif filename.find('test')==0:
            te_path=cur_path

print(f'training file: {tr_path}')
print(f'testing file: {te_path}')
```

```
../dataset/train.csv
../dataset/test.csv
training file: ../dataset/train.csv
testing file: ../dataset/test.csv
```

1.1 Preprocessing and Data Analysis

```
[13]: tr_df=pd.read_csv(tr_path)
tr_df.head()
```

```
[13]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
[14]: te_df=pd.read_csv(te_path)
te_df.head()
```

```
[14]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001015	Male	Yes	0	Graduate	No	
1	LP001022	Male	Yes	1	Graduate	No	
2	LP001031	Male	Yes	2	Graduate	No	
3	LP001035	Male	Yes	2	Graduate	No	
4	LP001051	Male	No	0	Not Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5720	0	110.0	360.0	
1	3076	1500	126.0	360.0	
2	5000	1800	208.0	360.0	
3	2340	2546	100.0	360.0	
4	3276	0	78.0	360.0	

	Credit_History	Property_Area
0	1.0	Urban
1	1.0	Urban
2	1.0	Urban
3	NaN	Urban
4	1.0	Urban

```
[15]: print(f'training set size[{tr_df.shape}]')
print(f'testing set size[{te_df.shape}]')
```

training set size[(614, 13)]

testing set size[(367, 12)]

the Preprocessing of the training dataset

```
[19]: tr_df.info(verbose=True,show_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education             614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome       614 non-null   int64
7   CoapplicantIncome     614 non-null   float64
8   LoanAmount            592 non-null   float64
9   Loan_Amount_Term      600 non-null   float64
```

```

10 Credit_History      564 non-null    float64
11 Property_Area       614 non-null    object
12 Loan_Status         614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB

```

```
[20]: # summary statistics
tr_df.describe()
```

```
[20]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term \
count	614.000000	614.000000	592.000000	600.00000
mean	5403.459283	1621.245798	146.412162	342.00000
std	6109.041673	2926.248369	85.587325	65.12041
min	150.000000	0.000000	9.000000	12.00000
25%	2877.500000	0.000000	100.000000	360.00000
50%	3812.500000	1188.500000	128.000000	360.00000
75%	5795.000000	2297.250000	168.000000	360.00000
max	81000.000000	41667.000000	700.000000	480.00000

	Credit_History
count	564.000000
mean	0.842199
std	0.364878
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

```
[24]: # the Id column is useless, drop it
if 'Loan_ID' in tr_df.columns:
    tr_df.drop('Loan_ID',axis=1,inplace=True)
if 'Loan_ID' in te_df.columns:
    te_df.drop('Loan_ID',axis=1,inplace=True)
# check the new shapes
print(f'training set[{tr_df.shape}]')
print(f'testing set[{te_df.shape}]')
```

```

training set[(614, 12)]
testing set[(367, 11)]

```

2 Missing values

填充数据集中的 Null value

```
[28]: tr_df.isnull().sum().sort_values(ascending=False)

null_cols=tr_df.isnull().any()
```

```
# Select only those entries that are True
null_cols=null_cols[null_cols].index.tolist()
print(null_cols)
```

```
['Gender', 'Married', 'Dependents', 'Self_Employed', 'LoanAmount',
'Loan_Amount_Term', 'Credit_History']
```

使用最常出现的值进行填充

```
[29]: # fill in the missing data
print('Before filling missing values\n\n', '#'*50, '\n')

for col in null_cols:
    print(f'{col}:\n{tr_df[col].value_counts()}\n', '-'*50)
    tr_df[col]=tr_df[col].fillna(tr_df[col].dropna().mode().values[0])

tr_df.isnull().sum().sort_values(ascending=False)
print('After filling missing values\n\n', '#'*50, '\n')
for col in null_cols:
    print(f'\n{col}:\n{tr_df[col].value_counts()}\n', '-'*50)
```

Before filling missing values

```
#####
```

Gender:

Gender

Male 489

Female 112

Name: count, dtype: int64

```
-----
```

Married:

Married

Yes 398

No 213

Name: count, dtype: int64

```
-----
```

Dependents:

Dependents

0 345

1 102

2 101

3+ 51

Name: count, dtype: int64

```
-----
```

Self_Employed:

Self_Employed

No 500

Yes 82

Name: count, dtype: int64

LoanAmount:

LoanAmount

120.0	20
110.0	17
100.0	15
160.0	12
187.0	12

..	
240.0	1
214.0	1
59.0	1
166.0	1
253.0	1

Name: count, Length: 203, dtype: int64

Loan_Amount_Term:

Loan_Amount_Term

360.0	512
180.0	44
480.0	15
300.0	13
240.0	4
84.0	4
120.0	3
60.0	2
36.0	2
12.0	1

Name: count, dtype: int64

Credit_History:

Credit_History

1.0	475
0.0	89

Name: count, dtype: int64

After filling missing values

#####

Gender:

Gender

Male	502
Female	112

Name: count, dtype: int64

```
Married:
Married
Yes      401
No       213
Name: count, dtype: int64
```

```
Dependents:
Dependents
0        360
1        102
2        101
3+        51
Name: count, dtype: int64
```

```
Self_Employed:
Self_Employed
No        532
Yes        82
Name: count, dtype: int64
```

```
LoanAmount:
LoanAmount
120.0      42
110.0      17
100.0      15
160.0      12
187.0      12
..
240.0       1
214.0       1
59.0        1
166.0       1
253.0       1
Name: count, Length: 203, dtype: int64
```

```
Loan_Amount_Term:
Loan_Amount_Term
360.0      526
180.0      44
480.0      15
300.0      13
240.0       4
84.0        4
```

```

120.0      3
60.0       2
36.0       2
12.0       1
Name: count, dtype: int64
-----

```

```

Credit_History:
Credit_History
1.0      525
0.0      89
Name: count, dtype: int64
-----

```

2.1 Data visualization

1. 将数据分类: categorical and numerical data

2.2 Loan status distribution

```

[32]: # list of all the numeric columns
num = tr_df.select_dtypes('number').columns.to_list()

# list of all the categorical columns
cat = tr_df.select_dtypes('object').columns.to_list()

print(f'number[{num}]')
print(f'categories[{cat}]')

# numeric df
loan_num=tr_df[num]
# categorical df
loan_cat=tr_df[cat]

```

```

number[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
'Loan_Amount_Term', 'Credit_History']]
categories[['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',
'Property_Area', 'Loan_Status']]

```

```

[35]: print(tr_df[cat[-1]].value_counts())
print(tr_df[cat[-1]])

total=float(len(tr_df[cat[-1]]))
plt.figure(figsize=(8,10))
sns.set(style='whitegrid')
ax=sns.countplot(data=tr_df,x=cat[-1])
for p in ax.patches:
    height=p.get_height()

```



```
ax.text(p.get_x()+p.get_width()/2.,height+3,'{:1.2f}'.format(height/  
↪total),ha='center')  
plt.show()
```

Loan_Status

Y 422

N 192

Name: count, dtype: int64

0 Y

1 N

2 Y

3 Y

4 Y

..

609 Y

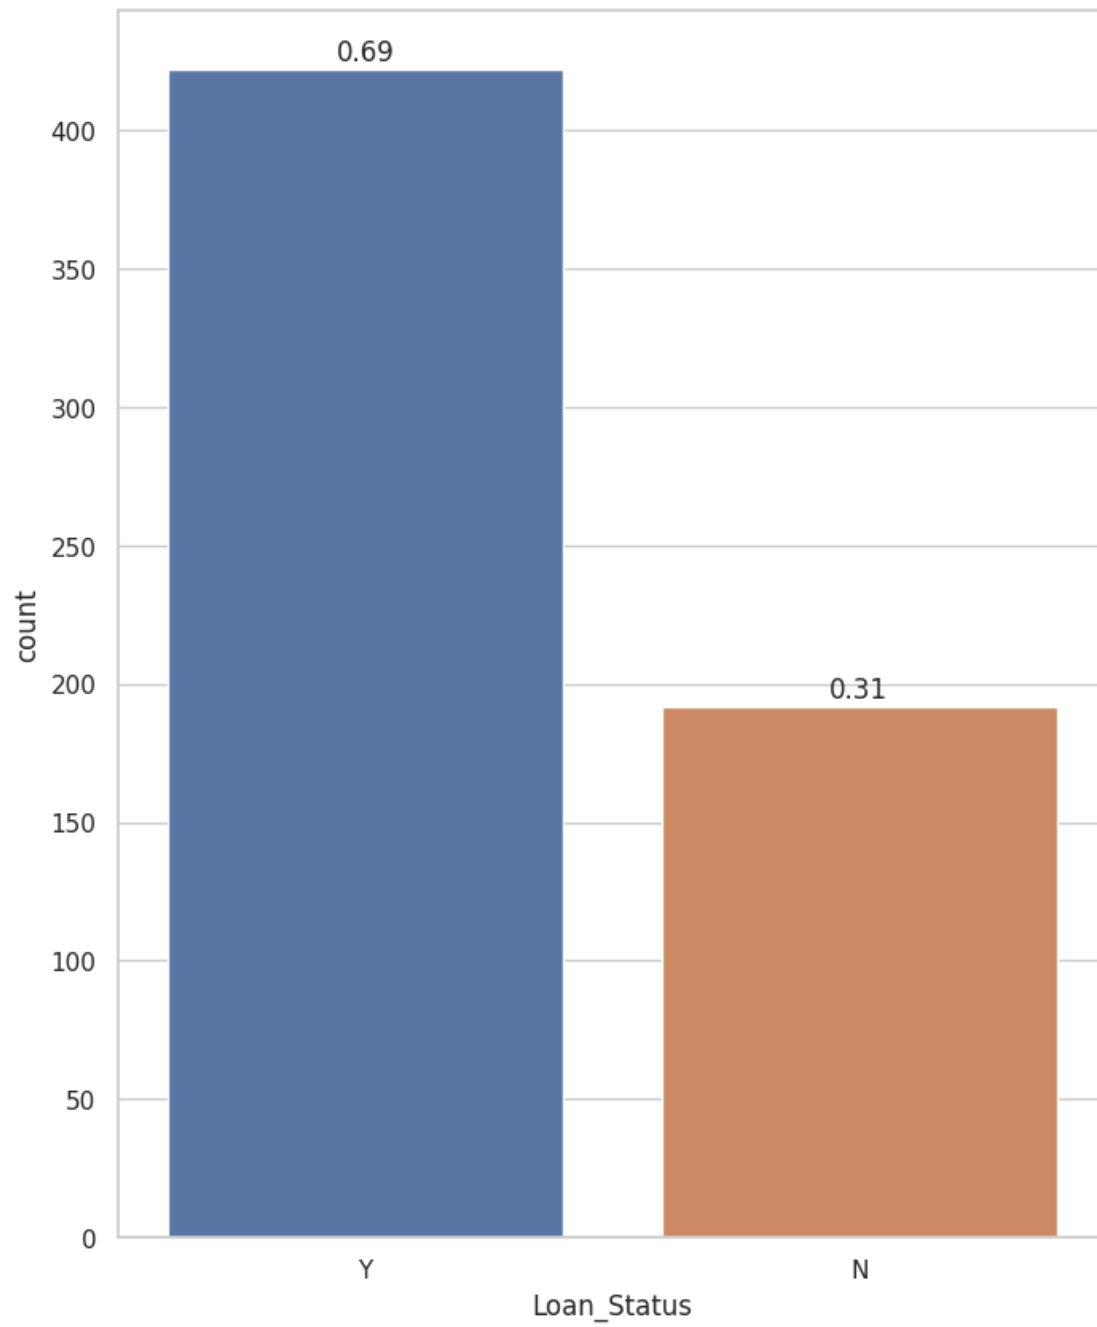
610 Y

611 Y

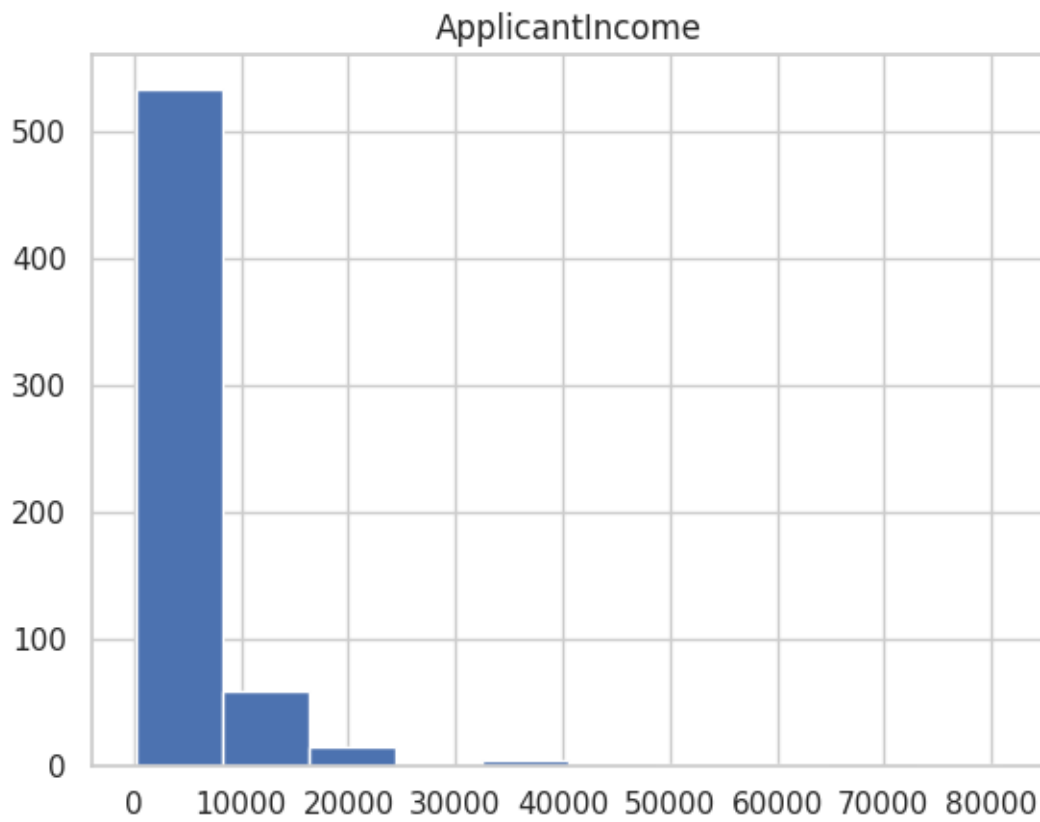
612 Y

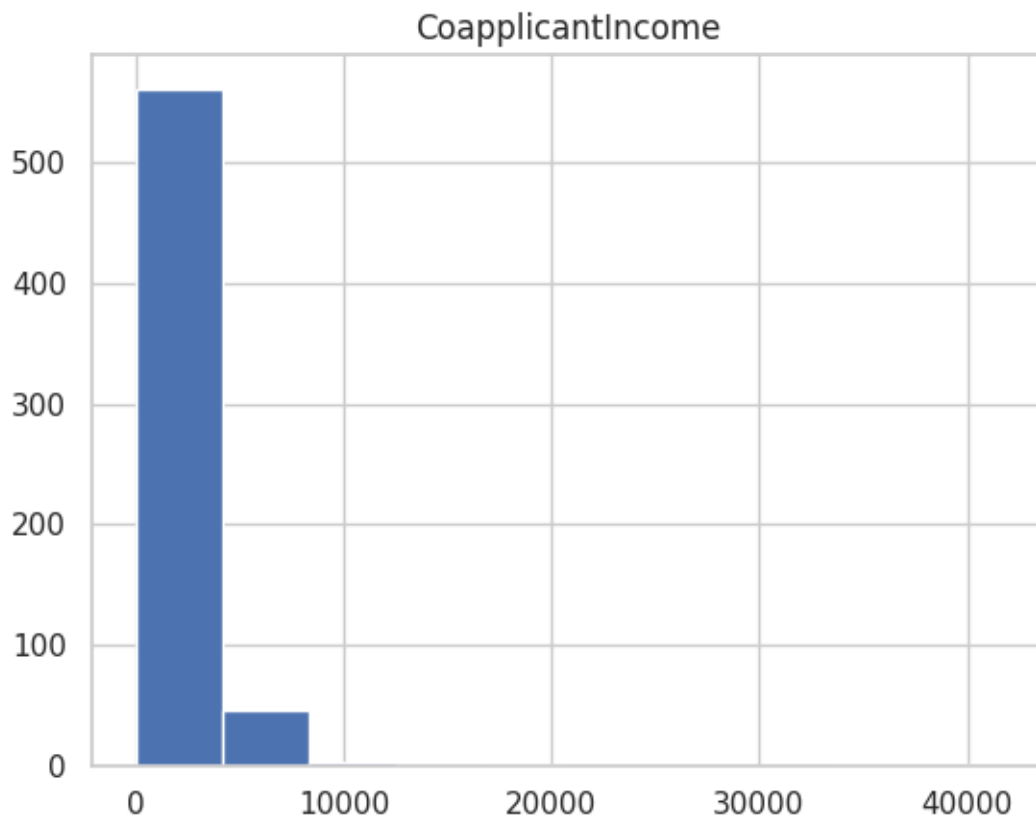
613 N

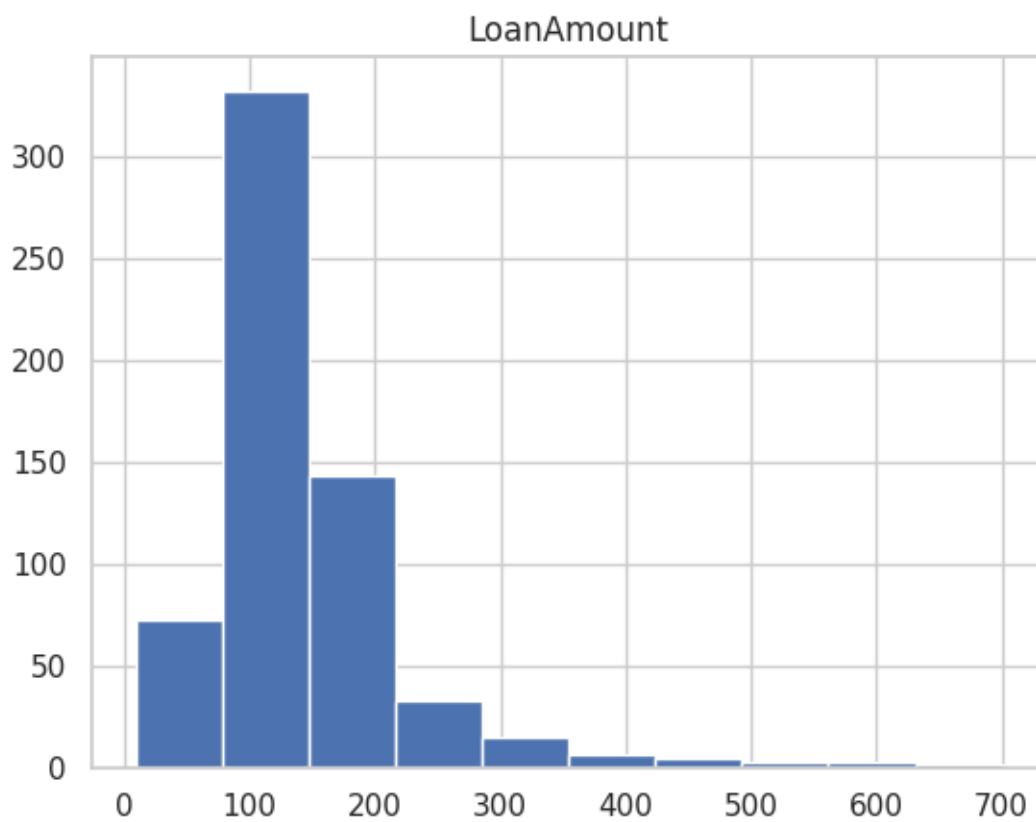
Name: Loan_Status, Length: 614, dtype: object

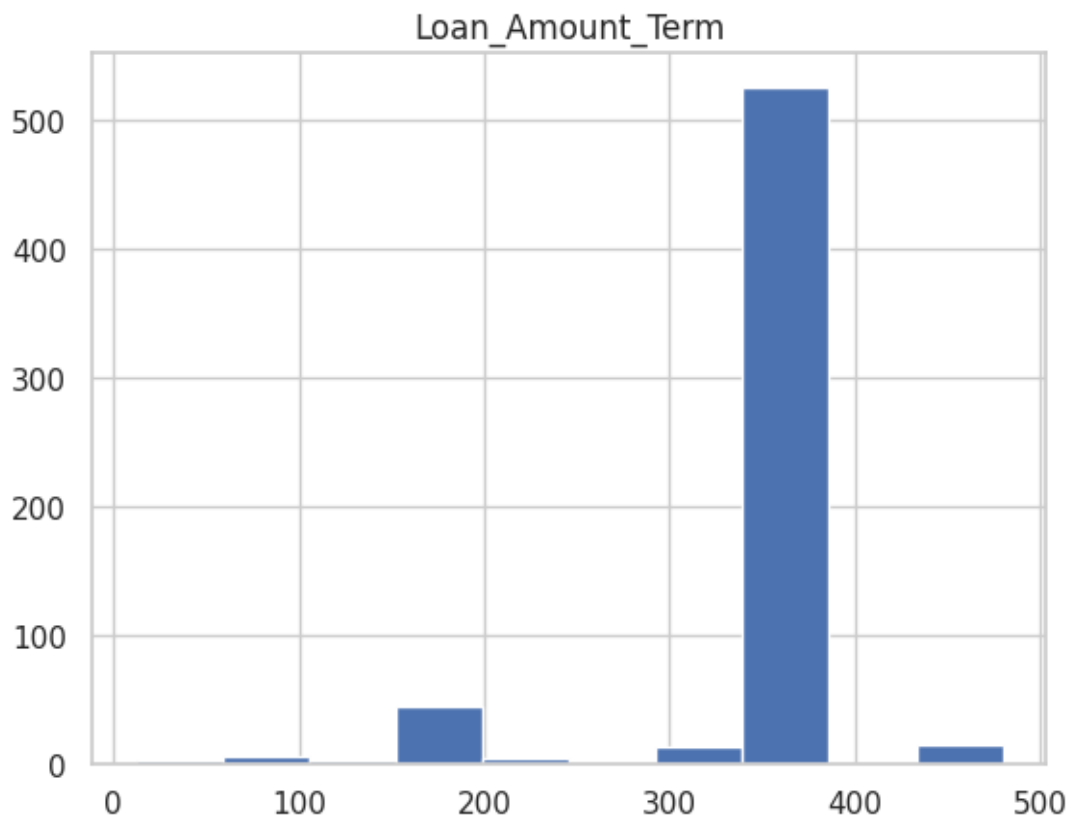


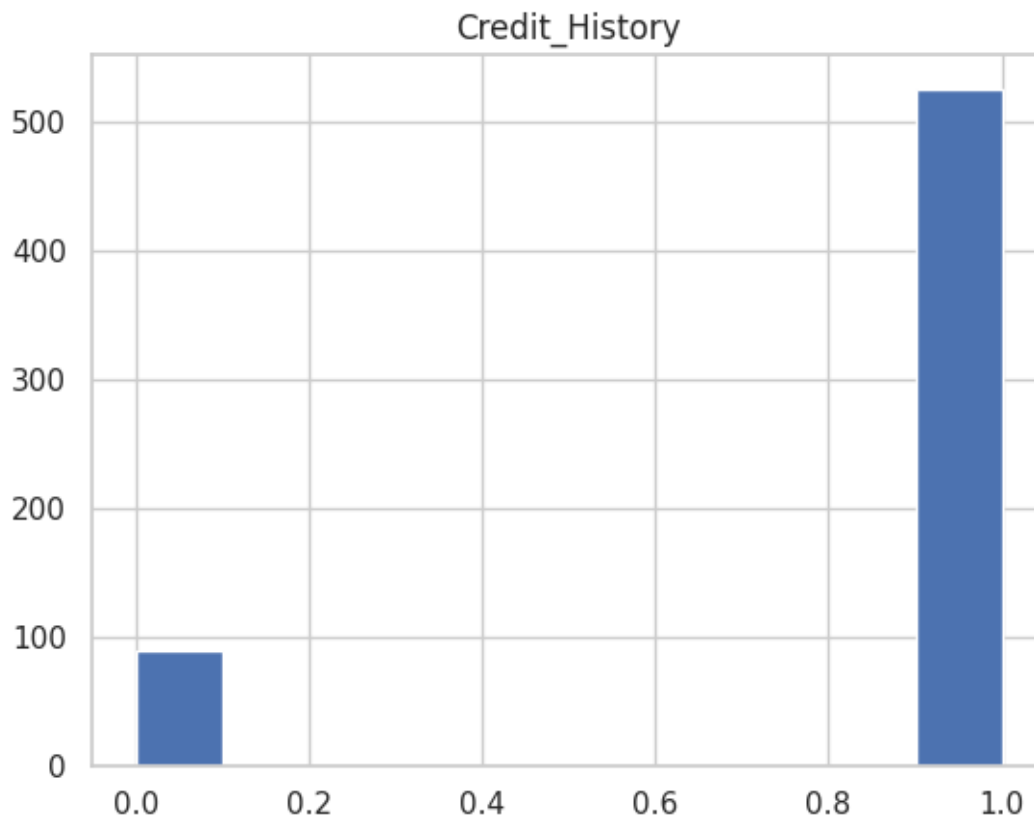
```
[37]: # Plot numeric columns
for i in loan_num:
    plt.hist(loan_num[i])
    plt.title(i)
    plt.show()
```












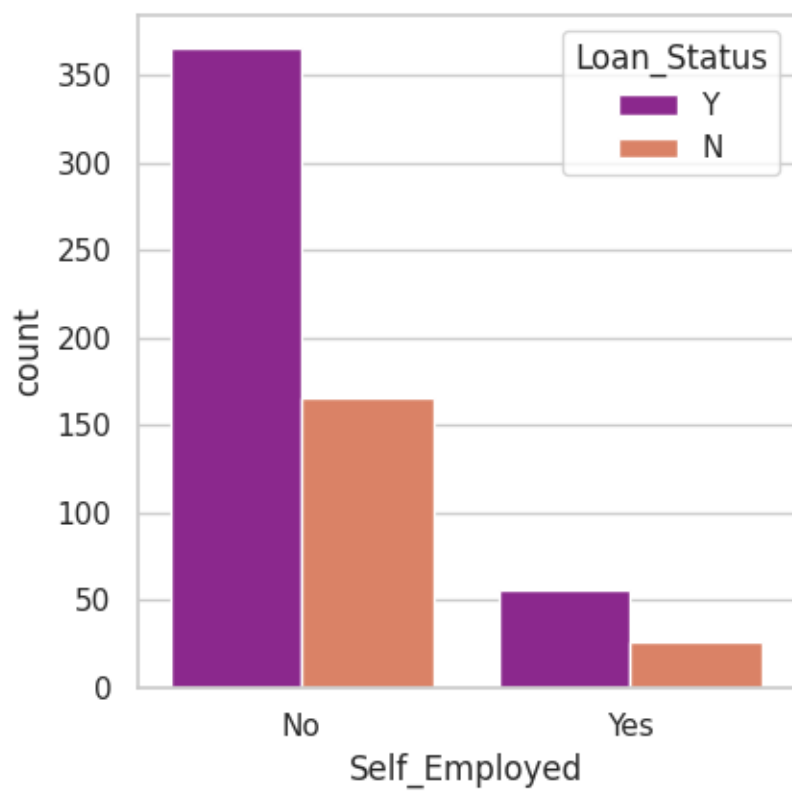


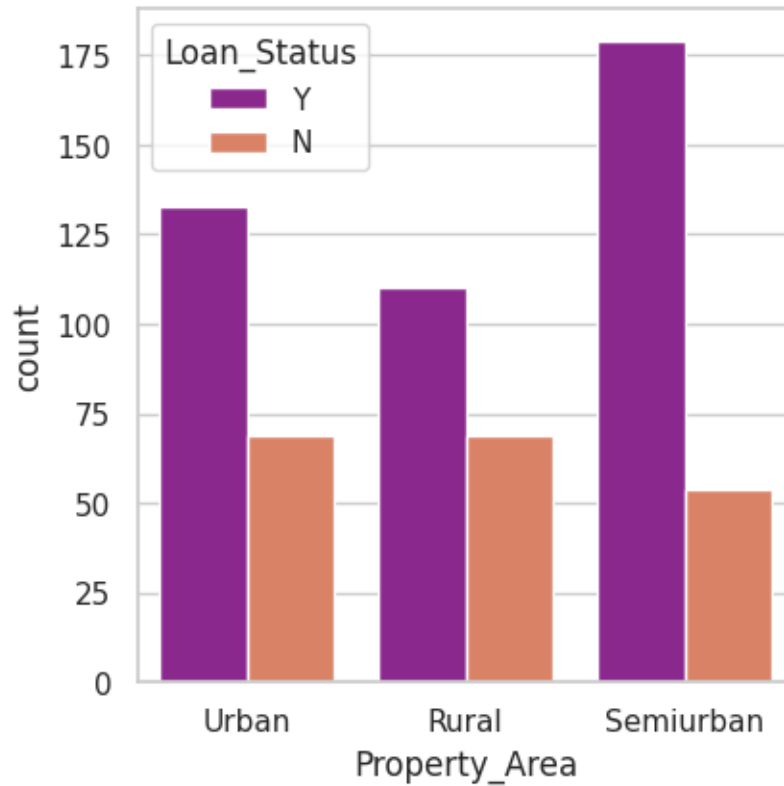
```
[40]: # Visualize categorical columns
for i in cat[:-1]:
    plt.figure(figsize=(15,10))
    plt.subplot(2,3,1)
    sns.countplot(x=i,hue='Loan_Status',data=tr_df,palette='plasma')
    plt.xlabel(i,fontsize=12)
```











2.3 Encoding data to numeric

```
[42]: # convert categorical values to numbers
to_numeric = {
    "Male": 1,
    "Female": 2,
    "Yes": 1,
    "No": 2,
    "Graduate": 1,
    "Not Graduate": 2,
    "Urban": 3,
    "Semiurban": 2,
    "Rural": 1,
    "Y": 1,
    "N": 0,
    "3+": 3,
}
# adding the new numeric values from the to_numeric variable to both datasets
tr_df = tr_df.applymap(lambda label: to_numeric.get(label) if label in to_numeric else label)
```

```

te_df=te_df.applymap(lambda label:to_numeric.get(label) if label in to_numeric_
↳else label)

# convert the dependents column
Dependents=pd.to_numeric(tr_df['Dependents'])
Dependents__=pd.to_numeric(te_df['Dependents'])

# drop the previous Dependents column
if 'Dependents' in tr_df.columns.to_list():
    tr_df.drop(['Dependents'],axis=1,inplace=True)
if 'Dependents' in te_df.columns.to_list():
    te_df.drop(['Dependents'],axis=1,inplace=True)

# concatenate the new Dependents column with both datasets
tr_df=pd.concat([tr_df,Dependents_],axis=1)
te_df=pd.concat([te_df,Dependents__],axis=1)

# check the dataset for validation
print(f'training set[{tr_df.shape}i], {tr_df.info()}')
print(f'testing set[{te_df.shape}], {te_df.info()}')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                614 non-null   int64
1   Married               614 non-null   int64
2   Education             614 non-null   int64
3   Self_Employed         614 non-null   int64
4   ApplicantIncome       614 non-null   int64
5   CoapplicantIncome     614 non-null   float64
6   LoanAmount            614 non-null   float64
7   Loan_Amount_Term      614 non-null   float64
8   Credit_History        614 non-null   float64
9   Property_Area         614 non-null   int64
10  Loan_Status           614 non-null   int64
11  Dependents            614 non-null   int64
dtypes: float64(4), int64(8)
memory usage: 57.7 KB
training set[(614, 12)i], None
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 367 entries, 0 to 366
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                356 non-null   float64
1   Married               367 non-null   int64

```

```

2   Education          367 non-null    int64
3   Self_Employed      344 non-null    float64
4   ApplicantIncome     367 non-null    int64
5   CoapplicantIncome   367 non-null    int64
6   LoanAmount          362 non-null    float64
7   Loan_Amount_Term    361 non-null    float64
8   Credit_History      338 non-null    float64
9   Property_Area       367 non-null    int64
10  Dependents          357 non-null    float64
dtypes: float64(6), int64(5)
memory usage: 31.7 KB
testing set[(367, 11)], None


```

2.4 Correlation matrix

To evaluate the relationship between each two categories in dataset

```
[43]: sns.heatmap(tr_df.corr(), cmap='cubehelix_r')
```

```
[43]: <Axes: >
```



```
[52]: # correlation table
corr=tr_df.corr()
corr.style.background_gradient(cmap='coolwarm').format(lambdax: '{:.2f}').
      format(x))
```

```
[52]: <pandas.io.formats.style.Styler at 0x7f81d5d3dcf0>
```

Credit_History has the highest correlation with Loan_Status, use this to predict target value

2.5 Logistic Regression

Divide dataset into two variables X as features y as Loan_Status the target value we want to predict

```
[53]: y=tr_df['Loan_Status']
X=tr_df.drop('Loan_Status',axis=1)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=0)
```

```
[54]: LR=LogisticRegression()
LR.fit(X_train,y_train)
y_hat=LR.predict(X_test)

# prediction summary by species
print(classification_report(y_test,y_hat))

# accuracy score
LR_SC=accuracy_score(y_hat,y_test)
print('accuracy is',accuracy_score(y_hat,y_test))
```

	precision	recall	f1-score	support
0	0.92	0.43	0.59	51
1	0.82	0.99	0.89	134
accuracy			0.83	185
macro avg	0.87	0.71	0.74	185
weighted avg	0.85	0.83	0.81	185

accuracy is 0.8324324324324325

```
[ ]:
```