

7_titanic_machine_learning_from_disaster

吴清柳

July 24, 2023

1 泰坦尼克：从灾难中生存

使用机器学习方法预测从泰坦尼克灾难中生存下来的概率

1.1 第一步：定义问题

根据船上人员的性别, 年龄, 职业等信息, 设计一种算法来预测泰坦尼克号上的乘客的生存概率

1.2 第二步：获取数据

```
pip install --upgrade kaggle
```

如果本地没有数据, 则从 Kaggle 上进行下载. 需要[设置账号和 API 到 ~/.kaggle/kaggle.json](#)

```
kaggle competitions download -c titanic
```

1.3 第三步：数据预处理

对数据进行清理

```
[3]: import sys
import pandas as pd
import matplotlib
import numpy as np
import scipy as sp
import IPython
import sklearn

import random
import time

# ignore warnings
import warnings
warnings.filterwarnings('ignore')
print('-'*25)

# Input data files are available in the `../dataset/` directory
from subprocess import check_output
print(check_output(['ls', '../dataset']).decode('utf8'))
```

```
-----  
gender_submission.csv  
test.csv  
titanic.zip  
train.csv
```

1.4 加载数据建模库

使用 `scikit-learn` 库来开发机器学习算法. 在 `sklearn` 中, 算法被叫做 `Estimators`, 实现在他们各自的类里. 对于数据可视化, 使用 `matplotlib` 和 `seaborn` 库.

```
[8]: # common model algorithms  
from sklearn import svm, tree, linear_model, neighbors, naive_bayes, ensemble,  
    ↪discriminant_analysis, gaussian_process  
from xgboost import XGBClassifier  
  
# common model helpers  
from sklearn.preprocessing import OneHotEncoder, LabelEncoder  
from sklearn import feature_selection, model_selection, metrics  
  
# visualization  
import matplotlib as mpl  
import matplotlib.pyplot as plt  
import matplotlib.pylab as pylab  
import seaborn as sns  
from pandas.plotting import scatter_matrix  
  
# configure visualization defaults  
# show plots in jupyter notebook inplace  
%matplotlib inline  
mpl.style.use('ggplot')  
sns.set_style('white')  
pylab.rcParams['figure.figsize']=12,8
```

1.5 了解数据

了解数据的形状 (数据类型, 数据值) 等等

1. Survived 代表乘客是否存活
2. PassengerID 和 Ticket 被假设为随机独立标识符, 对输出没有影响, 因此会被从分析中移除
3. Pclass 代表票型, 并映射社会经济状态, 表示 1 = 上层阶级, 2 = 中层阶级, 3 = 下层阶级;
4. Name 是名字数据类型, 可能可以在特征工程中根据 title 判断性别, 从 surname 中家庭大小.
5. Sex 和 Embarked 变量是命名数据类型. 会被转为 dummy 变量来进行数学计算.
6. Age 和 Fare 变量是连续量化数据类型;
7. SibSp 表示同在船上的兄弟姐妹的数量, Parch 表示同在船上的父母孩子. 都是离散量化数据类型. 可以在特征工程中建立一个家庭大小, 是孤立的变量;
8. Cabin 变量是命名数据类型, 可以在特征工程中大致定位事故发生时在船上的位置, 以及根据等级判断接机. 然而, 犹豫有许多 Null 值, 该变量用处不大, 被排除在分析之外;

```
[11]: # import data
data_raw = pd.read_csv("../dataset/train.csv")

# break dataset into train, test and validation
# the test file provided is the validation file for competition submission
# split the train set into train and test data
data_val = pd.read_csv("../dataset/test.csv")

# to play with our data, create a copy
# python assignment or equal passes by reference vs values
data1 = data_raw.copy(deep=True)

# however passing by reference is convenient, because learn both datasets at
# once
data_cleaner = [data1, data_val]

# preview data
print(data_raw.info())
data_raw.sample(10)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null   int64
1   Survived         891 non-null   int64
2   Pclass           891 non-null   int64
3   Name             891 non-null   object
4   Sex              891 non-null   object
5   Age              714 non-null   float64
6   SibSp            891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket           891 non-null   object
9   Fare             891 non-null   float64
10  Cabin            204 non-null   object
11  Embarked         889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```

```
[11]:
```

	PassengerId	Survived	Pclass	Name	Sex
838	839	1	3	Chip, Mr. Chang	male
585	586	1	1	Taussig, Miss. Ruth	female
201	202	0	3	Sage, Mr. Frederick	male
824	825	0	3	Panula, Master. Urho Abraham	male
864	865	0	2	Gill, Mr. John William	male
491	492	0	3	Windelov, Mr. Einar	male

646	647	0	3	Cor, Mr. Liudevit	male
690	691	1	1	Dick, Mr. Albert Adrian	male
97	98	1	1	Greenfield, Mr. William Bertram	male
592	593	0	3	Elsbury, Mr. William James	male

	Age	SibSp	Parch		Ticket	Fare	Cabin	Embarked
838	32.0	0	0		1601	56.4958	NaN	S
585	18.0	0	2		110413	79.6500	E68	S
201	NaN	8	2		CA. 2343	69.5500	NaN	S
824	2.0	4	1		3101295	39.6875	NaN	S
864	24.0	0	0		233866	13.0000	NaN	S
491	21.0	0	0	SOTON/OQ	3101317	7.2500	NaN	S
646	19.0	0	0		349231	7.8958	NaN	S
690	31.0	1	0		17474	57.0000	B20	S
97	23.0	0	1		PC 17759	63.3583	D10 D12	C
592	47.0	0	0		A/5 3902	7.2500	NaN	S

1.5.1 数据清理: 数据纠正, 数据补全, 数据创建和数据转换

在该阶段中, 1) 修正不正常数据和离群数据, 2) 完成丢失的信息, 3) 为分析创建新的特征, 4) 转换数据到正确的格式以用于计算和表示

```
[12]: print('Train columns with null values:\n', data1.isnull().sum())
      print('-'*10)

      print('Test/Validation columns with null values:\n', data_val.isnull().sum())
      print('-'*10)

      data_raw.describe(include='all')
```

Train columns with null values:

```
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64
```

Test/Validation columns with null values:

```
PassengerId      0
Pclass           0
```

```

Name          0
Sex           0
Age          86
SibSp        0
Parch        0
Ticket       0
Fare         1
Cabin       327
Embarked     0
dtype: int64
-----

```

```

[12]:      PassengerId  Survived  Pclass      Name  Sex  \
count      891.000000    891.000000    891.000000      891    891
unique         NaN         NaN         NaN      891      2
top           NaN         NaN         NaN  Braund, Mr. Owen Harris  male
freq          NaN         NaN         NaN          1    577
mean      446.000000     0.383838     2.308642         NaN    NaN
std       257.353842     0.486592     0.836071         NaN    NaN
min        1.000000     0.000000     1.000000         NaN    NaN
25%       223.500000     0.000000     2.000000         NaN    NaN
50%       446.000000     0.000000     3.000000         NaN    NaN
75%       668.500000     1.000000     3.000000         NaN    NaN
max       891.000000     1.000000     3.000000         NaN    NaN

```

```

      Age  SibSp  Parch  Ticket  Fare  Cabin  \
count  714.000000    891.000000    891.000000      891    891.000000      204
unique         NaN         NaN         NaN      681         NaN      147
top           NaN         NaN         NaN    347082         NaN    B96 B98
freq          NaN         NaN         NaN          7         NaN          4
mean     29.699118     0.523008     0.381594         NaN    32.204208         NaN
std     14.526497     1.102743     0.806057         NaN    49.693429         NaN
min       0.420000     0.000000     0.000000         NaN     0.000000         NaN
25%     20.125000     0.000000     0.000000         NaN     7.910400         NaN
50%     28.000000     0.000000     0.000000         NaN    14.454200         NaN
75%     38.000000     1.000000     0.000000         NaN    31.000000         NaN
max     80.000000     8.000000     6.000000         NaN   512.329200         NaN

```

```

      Embarked
count      889
unique        3
top          S
freq       644
mean       NaN
std       NaN
min       NaN
25%       NaN

```

50%	NaN
75%	NaN
max	NaN

```
[13]: # completing: complete or delete missing values in train and test/validation
# dataset
for dataset in data_cleaner:
    # complete missing age with median
    dataset['Age'].fillna(dataset['Age'].median(), inplace=True)

    # complete embarked with mode
    dataset['Embarked'].fillna(dataset['Embarked'].mode()[0], inplace=True)

    # complete missing fare with median
    dataset['Fare'].fillna(dataset['Fare'].median(), inplace=True)

# delete the cabin feature/column and others previously stated to exclude in
# train dataset
drop_column=['PassengerId','Cabin','Ticket']
data1.drop(drop_column,axis=1,inplace=True)

print(data1.isnull().sum())
print('-'*10)
print(data_val.isnull().sum())
```

```
Survived    0
Pclass      0
Name        0
Sex         0
Age         0
SibSp       0
Parch       0
Fare        0
Embarked    0
dtype: int64
-----
PassengerId    0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin          327
Embarked       0
dtype: int64
```

```

[15]: # create: feature engineering for train and test/validation dataset
for dataset in data_cleaner:
    # discrete variables
    dataset["FamilySize"] = dataset["SibSp"] + dataset["Parch"] + 1

    dataset["IsAlone"] = 1 # initialize to yes/1 is alone
    dataset["IsAlone"].loc[
        dataset["FamilySize"] > 1
    ] = 0 # now update to no/0 if family size is greater than 1

    # quick and dirty code split title from name
    dataset["Title"] = (
        dataset["Name"]
        .str.split(", ", expand=True)[1]
        .str.split(".", expand=True)[0]
    )

    # continuous variable bins using qcut, cut into bins with approximately
    # equal amount of elements
    dataset["FareBin"] = pd.qcut(dataset["Fare"], 4)

    # age bins/buckets using cut with the same interval of each bin
    dataset["AgeBin"] = pd.cut(dataset["Age"].astype(int), 5)

    # cleanup rare title names
    print(data1["Title"].value_counts())
    stat_min = (
        10 # while small is arbitrary, we'll use the common minimum in statistics
    )
    title_names = (
        data1["Title"].value_counts() < stat_min
    ) # this will create a true false series with title name as index

    # apply and lambda functions are quick and dirty code to find and replace with
    # fewer lines of code
    data1["Title"] = data1["Title"].apply(
        lambda x: "Misc" if title_names.loc[x] == True else x
    )
    print(data1['Title'].value_counts())
    print('-'*10)

    # preview data again
    data1.info()
    data_val.info()
    data1.sample(10)

```

Title

Mr

517

```

Miss          182
Mrs           125
Master        40
Dr            7
Rev           6
Mlle          2
Major         2
Col           2
the Countess  1
Capt         1
Ms            1
Sir           1
Lady          1
Mme           1
Don           1
Jonkheer      1

```

Name: count, dtype: int64

Title

```

Mr           517
Miss         182
Mrs          125
Master       40
Misc         27

```

Name: count, dtype: int64

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 891 entries, 0 to 890

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Survived	891 non-null	int64
1	Pclass	891 non-null	int64
2	Name	891 non-null	object
3	Sex	891 non-null	object
4	Age	891 non-null	float64
5	SibSp	891 non-null	int64
6	Parch	891 non-null	int64
7	Fare	891 non-null	float64
8	Embarked	891 non-null	object
9	FamilySize	891 non-null	int64
10	IsAlone	891 non-null	int64
11	Title	891 non-null	object
12	FareBin	891 non-null	category
13	AgeBin	891 non-null	category

dtypes: category(2), float64(2), int64(6), object(4)

memory usage: 85.9+ KB

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 418 entries, 0 to 417

Data columns (total 16 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	418 non-null	int64
1	Pclass	418 non-null	int64
2	Name	418 non-null	object
3	Sex	418 non-null	object
4	Age	418 non-null	float64
5	SibSp	418 non-null	int64
6	Parch	418 non-null	int64
7	Ticket	418 non-null	object
8	Fare	418 non-null	float64
9	Cabin	91 non-null	object
10	Embarked	418 non-null	object
11	FamilySize	418 non-null	int64
12	IsAlone	418 non-null	int64
13	Title	418 non-null	object
14	FareBin	418 non-null	category
15	AgeBin	418 non-null	category

dtypes: category(2), float64(2), int64(6), object(6)

memory usage: 47.1+ KB

```
[15]:
```

	Survived	Pclass	Name \
321	0	3	Danoff, Mr. Yoto
847	0	3	Markoff, Mr. Marin
543	1	2	Beane, Mr. Edward
731	0	3	Hassan, Mr. Houssein G N
320	0	3	Dennis, Mr. Samuel
369	1	1	Aubart, Mme. Leontine Pauline
270	0	1	Cairns, Mr. Alexander
746	0	3	Abbott, Mr. Rossmore Edward
599	1	1	Duff Gordon, Sir. Cosmo Edmund ("Mr Morgan")
820	1	1	Hays, Mrs. Charles Melville (Clara Jennings Gr...

	Sex	Age	SibSp	Parch	Fare	Embarked	FamilySize	IsAlone	Title \
321	male	27.0	0	0	7.8958	S	1	1	Mr
847	male	35.0	0	0	7.8958	C	1	1	Mr
543	male	32.0	1	0	26.0000	S	2	0	Mr
731	male	11.0	0	0	18.7875	C	1	1	Mr
320	male	22.0	0	0	7.2500	S	1	1	Mr
369	female	24.0	0	0	69.3000	C	1	1	Misc
270	male	28.0	0	0	31.0000	S	1	1	Mr
746	male	16.0	1	1	20.2500	S	3	0	Mr
599	male	49.0	1	0	56.9292	C	2	0	Misc
820	female	52.0	1	1	93.5000	S	3	0	Mrs

FareBin AgeBin

321	(-0.001, 7.91]	(16.0, 32.0]
847	(-0.001, 7.91]	(32.0, 48.0]
543	(14.454, 31.0]	(16.0, 32.0]
731	(14.454, 31.0]	(-0.08, 16.0]
320	(-0.001, 7.91]	(16.0, 32.0]
369	(31.0, 512.329]	(16.0, 32.0]
270	(14.454, 31.0]	(16.0, 32.0]
746	(14.454, 31.0]	(-0.08, 16.0]
599	(31.0, 512.329]	(48.0, 64.0]
820	(31.0, 512.329]	(48.0, 64.0]

1.5.2 转换格式

将类别数据转换成 dummy 变量, 用于数学分析.

此外, 为数据建模定义 x(independent/features/explanatory/predictor/etc.) 和 y(dependent/target/outcome/response/etc.) 变量

```
[ ]: # convert: convert objects to category using Label Encoder for train and
# test/validation dataset

# code categorical data
label = LabelEncoder()
for dataset in data_cleaner:
    dataset["Sex_Code"] = label.fit_transform(dataset["Sex"])
    dataset["Embarked_Code"] = label.fit_transform(dataset["Embarked"])
    dataset["Title_Code"] = label.fit_transform(dataset["Title"])
    dataset["FareBin_Code"] = label.fit_transform(dataset["FareBin"])
    dataset["AgeBin_Code"] = label.fit_transform(dataset["AgeBin"])

# define y variable aka target/outcome
Target = ["Survived"]

# define x variables for original features aka feature selection
data1_x = [
    "Sex",
    "Pclass",
    "Embarked",
    "Title",
    "SibSp",
    "Parch",
    "Age",
    "Fare",
    "FamilySize",
    "IsAlone",
]
# pretty name/values for charts
data1_x_calc = [
    "Sex_Code",
```

```

    "Pclass",
    "Embarked_Code",
    "Title_Code",
    "SibSp",
    "Parch",
    "Age",
    "Fare",
] # code for algorithm calculation

# define x variables original w/bin features to remove continuous variables
data1_x_bin = [
    "Sex_Code",
    "Pclass",
    "Embarked_Code",
    "Title_Code",
    "FamilySize",
    "AgeBin_Code",
    "FareBin_Code",
]
data1_xy_bin = Target + data1_x_bin
print("Bin X Y: ", data1_xy_bin, "\n")

# define x and y variables for dummy features original
data1_dummy=pd.get_dummies(data1[data1_x])
data1_x_dummy

```